

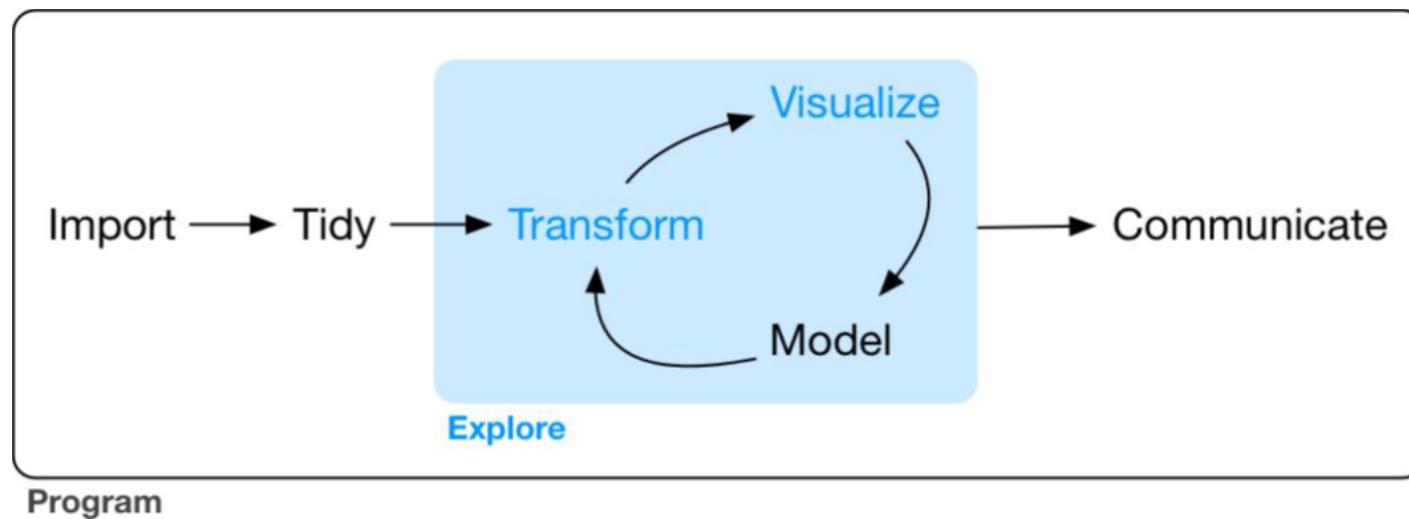
IE6600 Computation and Visualization for Analytics

Data Transformation with dplyr

(updated: 2022-06-13)

Data Transformation

Goal



Wickham, Hadley, and Garrett Grolemund. R For Data Science. OReilly, 2017.

Tips

Try press Alt-Shift-K.

Data Transformation with dplyr

Often you'll need to create some new variables or summaries, or maybe you just want to rename the variables or reorder the observations in order to make the data a little easier to work with.

```
library(nycflights13)
library(dplyr)
#or
library(tidyverse)
```

```
-- Attaching packages ----- tidyverse 1.2.1 --
v ggplot2 3.1.0      v purrr  0.3.0
v tibble  2.0.1      v dplyr  0.8.0.1
v tidyr   0.8.2      v stringr 1.4.0
v readr   1.3.1      v forcats 0.4.0
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
```

The conflict message tells you some of the other functions have been overwritten by Tidyverse. If you want to use the base version of these functions after loading dplyr, you'll need to use their full names: `stats::filter()` and `stats::lag()`, etc.

search()

```
search()
```

```
## [1] ".GlobalEnv"          "package:forcats"      "package:stringr"
## [4] "package:dplyr"        "package:purrr"        "package:readr"
## [7] "package:tidyr"        "package:tibble"       "package:ggplot2"
## [10] "package:tidyverse"    "package:nycflights13" "package:stats"
## [13] "package:graphics"     "package:grDevices"    "package:utils"
## [16] "package:datasets"     "package:methods"      "Autoloads"
## [19] "package:base"
```

You may use the full syntax `package::function_name()` to load the specific function, if there are any overwriting issues occurred.

nycflights13

To explore the basic data manipulation verbs of dplyr, we'll use `nycflights13::flights`. This data frame contains all 336,776 flights that departed from New York City in 2013:

```
head(flights)
```

```
## # A tibble: 6 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Interview data with view()

```
view(flights)
```


nycflights13 (cont'd)

You might also have noticed the row of three- (or four-) letter abbreviations under the column names. These describe the type of each variable:

- `int` stands for integers.
- `dbl` stands for doubles, or real numbers.
- `chr` stands for character vectors, or strings.
- `dtm` stands for date-times (a date + a time).

Flights actually is a tibble, a special type of data.frame. We will talk about it later.

nycflights13 (cont'd)

There are three other common types of variables that aren't used in this dataset but you'll encounter later.

- **lgl** stands for logical, vectors that contain only TRUE or FALSE.
- **fctr** stands for factors, which R uses to represent categorical variables with fixed possible values.
- **date** stands for dates.

dplyr Basics

In this slides you are going to learn the five key dplyr functions that allow you to solve the vast majority of your data-manipulation challenges:

- `filter()` pick observations by their values .
- `arrange()` Reorder the rows.
- `select()` Pick variables by their names.
- `mutate()` Create new variables with functions of existing variables.
- `summarize()` Collapse many values down to a single summary.
- `group_by()` Conjunction.

dplyr grammar

All verbs work similarly: `filter(df, argument,...)`

1. The first argument is a data frame.
2. The subsequent arguments describe what to do with the data frame, using the variable names (without quotes).
3. The result is a new data frame.

filter()

Filter rows with filter()

filter() allows you to subset observations based on their values.

Base function in R:

```
flights[flights$month==1&flights$day==1,]
```

filter() function in dplyr:

```
filter(flights, month==1, day==1)
```

Filter rows with filter() (cont'd)

We only want to see the Jan.1st flights

```
jan <- filter(flights, month==1, day==1)
jan
```

```
## # A tibble: 842 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # ... with 832 more rows, and 11 more variables: arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Filter rows with filter(): examples

```
dec25 <- filter(flights, month == 12, day == 25)
dec25
```

```
## # A tibble: 719 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013    12    25     456           500         -4     649           651
## 2  2013    12    25     524           515          9     805           814
## 3  2013    12    25     542           540          2     832           850
## 4  2013    12    25     546           550         -4    1022          1027
## 5  2013    12    25     556           600         -4     730           745
## 6  2013    12    25     557           600         -3     743           752
## 7  2013    12    25     557           600         -3     818           831
## 8  2013    12    25     559           600         -1     855           856
## 9  2013    12    25     559           600         -1     849           855
## 10 2013    12    25     600           600          0     850           846
## # ... with 709 more rows, and 11 more variables: arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```


filter - comparison

To use filtering effectively, you have to know how to select the observations that you want using the comparison operators. R provides the standard suite: >, >=, <, <=, != (not equal), and == (equal).

Flights on Feb~Dec, and before 28 th

```
filter(flights, month>1&!day>28)
```

Flights on Feb~Dec, and before 28 th

```
flights28 <- filter(flights, month>1&!day>28)
flights28
```

```
## # A tibble: 285,972 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013    10     1     447           500         -13     614           648
## 2  2013    10     1     522           517          5     735           757
## 3  2013    10     1     536           545         -9     809           855
## 4  2013    10     1     539           545         -6     801           827
## 5  2013    10     1     539           545         -6     917           933
## 6  2013    10     1     544           550         -6     912           932
## 7  2013    10     1     549           600        -11     653           716
## 8  2013    10     1     550           600        -10     648           700
## 9  2013    10     1     550           600        -10     649           659
## 10 2013    10     1     551           600         -9     727           730
## # ... with 285,962 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

tips

When you're starting out with R, the easiest mistake to make is to use `=` instead of `==` when testing for equality. When this happens you'll get an informative error:

```
filter(flights, month = 1)
```

```
## Error in `filter()`:  
## ! We detected a named input.  
## i This usually means that you've used `=` instead of `==`.  
## i Did you mean `month == 1`?
```

tips (cont'd)

There's another common problem you might encounter when using `==`: floating-point numbers.

```
sqrt(2)^2==2
```

```
## [1] FALSE
```

```
1/49*49==1
```

```
## [1] FALSE
```

tips (cont'd)

Computers use finite precision arithmetic (they obviously can't store an infinite number of digits!) so remember that every number you see is an [approximation](#). Instead of relying on `==`, use `near()`:


```
near(sqrt(2) ^ 2, 2)
```

```
## [1] TRUE
```

```
near(1 / 49 * 49, 1)
```

```
## [1] TRUE
```

Logical Operators



Flights on May or June

```
mayJune <- filter(flights, month==5|month==6)  
mayJune$month %>% unique()
```

```
## [1] 5 6
```

%in%

A useful shorthand for this problem is `x %in% y`. This will select every row where `x` is one of the values in `y`. We could use it to rewrite the preceding code:

Retrieve the flights information on Jan, Feb, and Mar

```
filter(flights, month%in%c(1,2,3)) %>% head()
```


!= not equals to

Flights not on Feb

```
filter(flights, month!=2)
```

```
## # A tibble: 311,825 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
## 10 2013     1     1     558           600          -2     753           745
## # ... with 311,815 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

|, &, and ", "

Compare the following three code chunks

```
filter(flights, !(arr_delay > 120 | dep_delay > 120)) %>%  
  select(dep_delay) %>% head(3)
```

```
filter(flights, arr_delay <= 120, dep_delay <= 120) %>%  
  select(dep_delay) %>% head(3)
```

```
filter(flights, !arr_delay > 120 & !dep_delay > 120) %>%  
  select(dep_delay) %>% head(3)
```

```
## # A tibble: 3 × 1  
##   dep_delay  
##   <dbl>  
## 1         2  
## 2         4  
## 3         2
```

Missing Values

One important feature of R that can make comparison tricky is missing values, or NAs (“not availables”).

```
NA>5
```

```
## [1] NA
```

```
NA==10
```

```
## [1] NA
```

```
NA+10
```

```
## [1] NA
```

```
NA/2
```

```
## [1] NA
```

NA==NA

The most confusing result is this one:

```
NA==NA
```

```
## [1] NA
```

But we can understand it easily in one example:

```
# Let ZAge be Zhenyuan's age. We don't know how old he is.  
ZAge <- NA  
# Let TAge be Trump's age. We don't know how old he is. (don't google it)  
TAge <- NA  
# Are Zhenyuan and Trump the same age?  
ZAge == TAge
```

```
## [1] NA
```

```
# We don't know!
```

NA with filter()

filter() only includes rows where the condition is TRUE; it excludes both FALSE and NA values. If you want to preserve missing values, ask for them explicitly:

```
df <- tibble(x = c(1, NA, 3))  
filter(df, x > 1)
```

```
## # A tibble: 1 × 1  
##       x  
##   <dbl>  
## 1     3
```

```
filter(df, is.na(x) | x > 1)
```

```
## # A tibble: 2 × 1  
##       x  
##   <dbl>  
## 1    NA  
## 2     3
```

Some other functions

```
is.na(NA)
```

```
## [1] TRUE
```

```
df <- data.frame(A=c(1,NA,2))
```

```
na.omit(df)
```

```
##    A
```

```
## 1 1
```

```
## 3 2
```

```
sum(df[,1], na.rm=T)
```

```
## [1] 3
```

Exercise 1

Find all flights that:

1. Flew to Houston (IAH or HOU)
2. Were operated by United(UA), American(AA), or Delta(DL)
3. Departed in summer (July, August, and September)

Exercise 2

Use data set `msleep`, and create a new data frame of mammals with feeding type `carnivore` and brain weight less than the average of brain weight over all mammals. Make sure no NA values in column of brain weight.

arrange()

Arrange Rows with arrange()

arrange() works similarly to filter() except that instead of selecting rows, it changes their order

Base function in R:

```
flights[order(flights$year, flights$month, flights$day, decreasing=F),]
```

arrange() function in dplyr:

```
arrange(flights, year, month, day)
```

Arrange Rows with arrange() (cont'd)

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
## 10 2013     1     1     558             600          -2     753             745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Use desc() to reorder by a column in descending order

```
arrange(flights, desc(arr_delay))
```

```
## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     9     641           900       1301    1242         1530
## 2  2013     6    15    1432          1935       1137    1607         2120
## 3  2013     1    10    1121          1635       1126    1239         1810
## 4  2013     9    20    1139          1845       1014    1457         2210
## 5  2013     7    22     845          1600       1005    1044         1815
## 6  2013     4    10    1100          1900        960    1342         2211
## 7  2013     3    17    2321           810        911     135         1020
## 8  2013     7    22    2257           759        898     121         1026
## 9  2013    12     5     756          1700        896    1058         2020
## 10 2013     5     3    1133          2055        878    1250         2215
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Missing values are always sorted at the end

```
df <- data.frame(x = c(5, 2, NA))  
arrange(df, x) #or arrange(df, desc(x))
```

```
##      x  
## 1    2  
## 2    5  
## 3  NA
```

Exercises

1. Sort flights to find the most delayed flights. Find the flights that left earliest.
2. Sort flights to find the fastest flights.

select()

Select Columns with select()

select() allows you to rapidly zoom in on a useful subset using operations based on the names of the variables.

Base function in R:

```
# Select columns by name  
flights[,c("year", "month", "day")]
```

select() function in dplyr:

```
# Select columns by name  
select(flights, year, month, day)
```


Select Columns with `select()` (cont'd)

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 × 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

Select all columns between year and day

```
select(flights, year:day)
```

```
## # A tibble: 336,776 × 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

Select all columns except those from year to day

```
select(flights, -(year:day))
```

```
## # A tibble: 336,776 × 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
##   <int>         <int>         <dbl>   <int>         <int>         <dbl> <chr>
## 1     517           515           2     830           819           11 UA
## 2     533           529           4     850           830           20 UA
## 3     542           540           2     923           850           33 AA
## 4     544           545          -1    1004          1022          -18 B6
## 5     554           600          -6     812           837          -25 DL
## 6     554           558          -4     740           728           12 UA
## 7     555           600          -5     913           854           19 B6
## 8     557           600          -3     709           723          -14 EV
## 9     557           600          -3     838           846           -8 B6
## 10    558           600          -2     753           745            8 AA
## # ... with 336,766 more rows, and 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Other arguments within select()

There are a number of helper functions you can use within select():

- `starts_with("abc")` matches names that begin with "abc".
- `ends_with("xyz")` matches names that end with "xyz".
- `contains("ijk")` matches names that contain "ijk".
- `matches("^a")` selects variables that match a regular expression. (check R4DS "regular expressions")
- `num_range("x", 1:3)` matches x1, x2, and x3.

select(): examples

```
abc.df <- data.frame(apple=c("b", "c"), an.orange=1:2, orange1=2:3)
abc.df
```

```
##   apple an.orange orange1
## 1    b          1         2
## 2    c          2         3
```

```
select(abc.df, starts_with("app"))
```

```
##   apple
## 1    b
## 2    c
```

```
select(abc.df, ends_with("ge"))
```

```
##   an.orange  
## 1         1  
## 2         2
```

```
select(abc.df, contains("pp"))
```

```
##   apple  
## 1     b  
## 2     c
```

```
select(abc.df, matches("^a"))
```

```
##   apple an.orange  
## 1      b         1  
## 2      c         2
```

```
select(abc.df, num_range("orange", 1))
```

```
##   orange1  
## 1       2  
## 2       3
```

mutate()

Add New Variables with mutate()

Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns. That's the job of `mutate()`.

Create a new data frame

```
temp <- data.frame(A=c(1:3),B=c(2:4))
```

Base function in R:

```
temp$C <- temp$A-temp$B
```

`mutate()` function in dplyr:

```
mutate(temp, C=A-B)
```

Add New Variables with mutate() (cont'd)

```
temp <- data.frame(A=c(1:3),B=c(2:4))  
mutate(temp, C=A-B)
```

```
##   A B  C  
## 1 1 2 -1  
## 2 2 3 -1  
## 3 3 4 -1
```

transmute() for only keeping the new variables

If you only want to keep the new variables, use transmute():

```
transmute(temp, C=A-B)
```

```
##      C  
## 1 -1  
## 2 -1  
## 3 -1
```

Example I

```
flights_delay <- select(flights,
                        year:day,
                        ends_with("delay"),
                        distance,
                        air_time)
mutate(flights_delay,
       gain = arr_delay - dep_delay,
       speed = distance / air_time * 60)
```

A tibble: 336,776 × 9

##	year	month	day	dep_delay	arr_delay	distance	air_time	gain	speed
##	<int>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1	2013	1	1	2	11	1400	227	9 370.
##	2	2013	1	1	4	20	1416	227	16 374.
##	3	2013	1	1	2	33	1089	160	31 408.
##	4	2013	1	1	-1	-18	1576	183	-17 517.
##	5	2013	1	1	-6	-25	762	116	-19 394.
##	6	2013	1	1	-4	12	719	150	16 288.
##	7	2013	1	1	-5	19	1065	158	24 404.
##	8	2013	1	1	-3	-14	229	53	-11 259.
##	9	2013	1	1	-3	-8	944	140	-5 405.
##	10	2013	1	1	-2	8	733	138	10 319.

... with 336,766 more rows

Example II

Note that you can refer to columns that you've just created:

```
mutate(  
  flights_delay,  
  gain = arr_delay - dep_delay, # New column  
  hours = air_time / 60, # New column  
  gain_per_hour = gain / hours  
)
```

```
## # A tibble: 336,776 × 10  
##   year month   day dep_delay arr_delay distance air_time  gain hours  
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>     <dbl> <dbl> <dbl>  
## 1  2013     1     1         2         11    1400      227     9  3.78  
## 2  2013     1     1         4         20    1416      227    16  3.78  
## 3  2013     1     1         2         33    1089      160    31  2.67  
## 4  2013     1     1        -1        -18    1576      183   -17  3.05  
## 5  2013     1     1        -6        -25     762      116   -19  1.93  
## 6  2013     1     1        -4         12     719      150    16  2.5  
## 7  2013     1     1        -5         19    1065      158    24  2.63  
## 8  2013     1     1        -3        -14     229       53   -11  0.883  
## 9  2013     1     1        -3         -8     944      140    -5  2.33  
## 10 2013     1     1        -2          8     733      138    10  2.3
```

Example III with transmute() for only keep the new variables

```
transmute(  
  flights,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours  
)
```

```
## # A tibble: 336,776 × 3  
##       gain hours gain_per_hour  
##   <dbl> <dbl>         <dbl>  
## 1      9 3.78           2.38  
## 2     16 3.78           4.23  
## 3     31 2.67          11.6  
## 4    -17 3.05          -5.57  
## 5    -19 1.93          -9.83  
## 6     16 2.5            6.4  
## 7     24 2.63           9.11  
## 8    -11 0.883         -12.5  
## 9     -5 2.33          -2.14  
## 10    10 2.3            4.35
```

mutate() with ifelse()

```
temp <- data.frame(A=c(1:3),B=c(2:4))  
temp
```

```
##   A B  
## 1 1 2  
## 2 2 3  
## 3 3 4
```

mutate() with ifelse() (cont'd)

In column B, We would like to replace all the values with the following pattern: if observation >3, then 1; if observation <=3, then 0

```
mutate(temp, B=ifelse(B>3, 1, 0))
```

```
##   A B
## 1 1 0
## 2 2 0
## 3 3 1
```


rename(), recode()

rename()

Rename variables by name

```
temp <- data.frame(A=c(1:3),B=c(2:4))  
temp
```

```
##   A B  
## 1 1 2  
## 2 2 3  
## 3 3 4
```

rename() vs names()

Rename column A as A1, column B as B1

```
names(temp) <- c("A1", "B1")
```

```
rename(temp, A1=A, B1=B)
```

```
##   A1 B1  
## 1  1  2  
## 2  2  3  
## 3  3  4
```

recode()

Recode values

```
temp1 <- data.frame(A=letters[1:3], B=c(1:3))  
temp1
```

```
##   A B  
## 1 a 1  
## 2 b 2  
## 3 c 3
```

```
recode(temp1$A, "b"="a")
```

```
## [1] "a" "a" "c"
```

```
mutate(temp1, A=recode(temp1$A, "b"="2"))
```

```
##   A B  
## 1 a 1  
## 2 2 2  
## 3 c 3
```

```
# temp1 %>% mutate(A=recode(A, "b"="2"))
```

```
mutate(temp1, B=recode(temp1$B, "2"="b"))
```

```
##   A    B  
## 1 a <NA>  
## 2 b    b  
## 3 c <NA>
```

According to the coercion rule, all the elements should be in the same data type within one variable/vector. since "b" is char data type, when 2 (integer) has been replaced with b, the other two integer 1 and 3 will be overwritten as NA.

Exercise

Let's create one data frame as follows:

```
##   A B  C
## 1 a 1  2
## 2 b 2 NA
## 3 c 3 13
## 4 d 4 56
## 5 e 5 NA
```

Then we would like to replace all the NAs in column C. If the value is NA, we replace it with the mean of column C, otherwise just keep the value.

```
data.frame(A=letters[1:5], B=c(1:5), C=c(2, NA, 13, 56, NA)) %>%
  mutate(C=ifelse(is.na(C), mean(C,na.rm = T),C))
```

summarize()

Summaries with summarize()

The last key verb is `summarize()`. It collapses a data frame to a single row:

```
summarize(flights, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 × 1  
##   delay  
##   <dbl>  
## 1  12.6
```


summarize() with group_by()

For example, if we applied exactly the same code to a data frame grouped by date, we get the average delay per month:

```
flights %>%  
  group_by(year, month) %>%  
  summarize(delay=mean(dep_delay, na.rm = T))
```

```
flights %>%  
  group_by(year, month) %>%  
  summarize(delay=mean(dep_delay, na.rm = T))
```

`summarise()` has grouped output by 'year'. You can override using the
`.groups` argument.

```
## # A tibble: 12 × 3  
## # Groups:   year [1]  
##   year month delay  
##   <int> <int> <dbl>  
## 1  2013     1  10.0  
## 2  2013     2  10.8  
## 3  2013     3  13.2  
## 4  2013     4  13.9  
## 5  2013     5  13.0  
## 6  2013     6  20.8  
## 7  2013     7  21.7  
## 8  2013     8  12.6  
## 9  2013     9   6.72  
## 10 2013    10   6.24  
## 11 2013    11   5.44  
## 12 2013    12  16.6
```

Missing values with summarize()

You may have wondered about the `na.rm` argument we used earlier. What happens if we don't set it?

```
flights %>%  
  group_by(year, month, day) %>%  
  summarize(mean = mean(dep_delay))
```

```
## `summarise()` has grouped output by 'year', 'month'. You can override using the  
## `.groups` argument.
```

```
## # A tibble: 365 × 4  
## # Groups:   year, month [12]  
##   year month   day mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1    NA  
## 2  2013     1     2    NA  
## 3  2013     1     3    NA  
## 4  2013     1     4    NA  
## 5  2013     1     5    NA  
## 6  2013     1     6    NA  
## 7  2013     1     7    NA  
## 8  2013     1     8    NA
```

Removing missing values with na.rm in summarize()

```
flights %>%  
  group_by(year, month, day) %>%  
  summarize(mean = mean(dep_delay, na.rm = TRUE))
```

`summarise()` has grouped output by 'year', 'month'. You can override using the
`.groups` argument.

```
## # A tibble: 365 × 4  
## # Groups:   year, month [12]  
##   year month   day mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1 11.5  
## 2  2013     1     2 13.9  
## 3  2013     1     3 11.0  
## 4  2013     1     4  8.95  
## 5  2013     1     5  5.73  
## 6  2013     1     6  7.15  
## 7  2013     1     7  5.42  
## 8  2013     1     8  2.55  
## 9  2013     1     9  2.28
```

When you use the dplyr verbs on a grouped data frame they'll be automatically applied "by group." For example, if we applied exactly the same code to a data frame grouped by date, we get the average delay per date: [By year, month, and date](#)

```
by_day <- group_by(flights, year, month, day)
summarize(by_day, delay = mean(dep_delay, na.rm = TRUE))
```

```
## `summarise()` has grouped output by 'year', 'month'. You can override using the
## `.groups` argument.
```

```
## # A tibble: 365 × 4
## # Groups:   year, month [12]
##   year month   day delay
##   <int> <int> <int> <dbl>
## 1  2013     1     1  11.5
## 2  2013     1     2  13.9
## 3  2013     1     3  11.0
## 4  2013     1     4   8.95
## 5  2013     1     5   5.73
## 6  2013     1     6   7.15
## 7  2013     1     7   5.42
## 8  2013     1     8   2.55
## 9  2013     1     9   2.28
## 10 2013     1    10   2.84
## # ... with 355 more rows
```

By month

```
by_month <- group_by(flights, month)
summarize(by_month, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 12 × 2
##   month delay
##   <int> <dbl>
## 1     1  10.0
## 2     2  10.8
## 3     3  13.2
## 4     4  13.9
## 5     5  13.0
## 6     6  20.8
## 7     7  21.7
## 8     8  12.6
## 9     9   6.72
## 10    10   6.24
## 11    11   5.44
## 12    12  16.6
```

By year

```
by_year <- group_by(flights, year)
summarize(by_year, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 × 2
##   year delay
##   <int> <dbl>
## 1  2013  12.6
```


Counts

Whenever you do any aggregation, it's always a good idea to include either a count `n()`, or a count of nonmissing values `sum(!is.na(x))`.

```
# Look at the planes (identified by their tail number) that have the  
# highest average delays:  
not_cancelled <- flights %>%  
  filter(!is.na(dep_delay), !is.na(arr_delay))  
  
delays <- not_cancelled %>%  
  group_by(tailnum) %>%  
  summarize(delay = mean(arr_delay))  
delays
```

```
## # A tibble: 4,037 × 2  
##   tailnum delay  
##   <chr>   <dbl>  
## 1 D942DN  31.5  
## 2 N0EGMQ   9.98  
## 3 N10156  12.7  
## 4 N102UW   2.94  
## 5 N103US  -6.93  
## 6 N104UW   1.80
```


Counts (cont'd)




Counts (cont'd)

We can get more insight if we draw a scatterplot of number of flights versus average delay:


```
not_cancelled %>%  
  group_by(tailnum) %>%  
  summarize(delay = mean(arr_delay, na.rm = TRUE),  
            n = n()) %>%  
  ggplot(mapping = aes(x = n, y = delay, color = delay < 0)) +  
  geom_point(alpha = 1 / 10)
```

Counts (cont'd)



Filter out all delay times less than and equals to 25 to see more details.

```
not_cancelled %>%  
  group_by(tailnum) %>%  
  summarize(delay = mean(arr_delay, na.rm = TRUE),  
            n = n()) %>%  
  filter(n > 25) %>%  
  ggplot(mapping = aes(x = n, y = delay, color = delay < 0)) +  
  geom_point(alpha = 1 / 10)
```



Useful Summary Functions

Measure of location: `mean()` and `median()`

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(# average delay:
    avg_delay1 = mean(arr_delay),
    # average positive delay:
    avg_delay2 = mean(arr_delay[arr_delay > 0]))
```

`summarise()` has grouped output by 'year', 'month'. You can override using the
`.groups` argument.

```
## # A tibble: 365 × 5
## # Groups:   year, month [12]
##   year month   day avg_delay1 avg_delay2
##   <int> <int> <int>      <dbl>      <dbl>
## 1  2013     1     1      12.7        32.5
## 2  2013     1     2      12.7        32.0
## 3  2013     1     3       5.73        27.7
## 4  2013     1     4      -1.93        28.3
## 5  2013     1     5      -1.53        22.6
## 6  2013     1     6       4.24        24.4
```

Useful Summary Functions (cont'd)

Measure of spread `sd()`

Some destinations more variable than to others

```
not_cancelled %>%  
  group_by(dest) %>%  
  summarize(distance_sd = sd(distance)) %>%  
  arrange(desc(distance_sd))
```

```
## # A tibble: 104 × 2  
##   dest distance_sd  
##   <chr>         <dbl>  
## 1 EGE          10.5  
## 2 SAN          10.4  
## 3 SFO          10.2  
## 4 HNL          10.0  
## 5 SEA           9.98  
## 6 LAS           9.91  
## 7 PDX           9.87  
## 8 PHX           9.86  
## 9 LAX           9.66  
## 10 IND          9.46  
## # ... with 94 more rows
```


Useful Summary Functions (cont'd)

Measures of rank `min(x)`, `max(x)`

```
# When do the first and last flights leave each day?  
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(first = min(dep_time),  
            last = max(dep_time))
```

```
## `summarise()` has grouped output by 'year', 'month'. You can override using the  
## `.groups` argument.
```

```
## # A tibble: 365 × 5  
## # Groups:   year, month [12]  
##   year month   day first  last  
##   <int> <int> <int> <int> <int>  
## 1  2013     1     1   517  2356  
## 2  2013     1     2    42  2354  
## 3  2013     1     3    32  2349  
## 4  2013     1     4    25  2358  
## 5  2013     1     5    14  2357  
## 6  2013     1     6    16  2355  
## 7  2013     1     7    49  2359
```

Useful Summary Functions (cont'd)

Measures of position `first(x)`, `nth(x, 2)`, `last(x)`

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(first_dep = first(dep_time),
            second_dep = nth(dep_time, 2))
```

`summarise()` has grouped output by 'year', 'month'. You can override using the
`.groups` argument.

```
## # A tibble: 365 × 5
## # Groups:   year, month [12]
##   year month   day first_dep second_dep
##   <int> <int> <int>      <int>      <int>
## 1  2013     1     1         517         533
## 2  2013     1     2          42         126
## 3  2013     1     3          32          50
## 4  2013     1     4          25         106
## 5  2013     1     5          14          37
## 6  2013     1     6          16         458
## 7  2013     1     7          49         454
## 8  2013     1     8         454         524
```

Useful Summary Functions (cont'd)

Distinct values, `n_distinct(x)`


```
# Which destinations have the most carriers?
not_cancelled %>%
  group_by(dest) %>%
  summarize(carriers = n_distinct(carrier)) %>%
  arrange(desc(carriers))
```

```
## # A tibble: 104 × 2
##   dest carriers
##   <chr>    <int>
## 1 ATL         7
## 2 BOS         7
## 3 CLT         7
## 4 ORD         7
## 5 TPA         7
## 6 AUS         6
## 7 DCA         6
## 8 DTW         6
## 9 IAD         6
## 10 MSP        6
## # ... with 94 more rows
```

Combining Multiple Operations with the Pipe %>%

Imagine that we want to explore the relationship between the distance and average delay for each location.

```
by_dest <- group_by(flights, dest)
delay <- summarize(
  by_dest,
  count = n(),
  dist = mean(distance, na.rm = TRUE),
  delay = mean(arr_delay, na.rm = TRUE)
)
delay <- filter(delay, count > 20, dest != "HNL")
```



There are three steps to prepare this data:

- Group flights by destination.
- Summarize to compute distance, average delay, and number of flights.
- Filter to remove noisy points and Honolulu airport, which is almost twice as far away as the next closest airport.

Combining Multiple Operations with the Pipe %>%

There's another way to tackle the same problem with the pipe, %>%:

```
flights %>%  
  group_by(dest) %>%  
  summarize(  
    count = n(),  
    dist = mean(distance, na.rm = TRUE),  
    delay = mean(arr_delay, na.rm = TRUE)  
  ) %>%  
  filter(count > 20, dest != "HNL") %>%  
  ggplot(mapping = aes(x = dist, y = delay)) +  
  geom_point(aes(size = count), alpha = 1 / 3) +  
  geom_smooth(se = FALSE)
```

Combining Multiple Operations with the Pipe %>%

There's another way to tackle the same problem with the pipe, %>%:

small data set: transmute() vs summarize()

```
a <-  
  data.frame(  
    c = c("111", "112", "111", "113"),  
    a = c("a", "b", "b", "c"),  
    b = 1:4  
  )  
a
```

```
##      c a b  
## 1 111 a 1  
## 2 112 b 2  
## 3 111 b 3  
## 4 113 c 4
```

transmute()

```
a %>%  
  group_by(c) %>% transmute(a1 = n_distinct(a))
```

```
## # A tibble: 4 × 2  
## # Groups:   c [3]  
##   c      a1  
##   <chr> <int>  
## 1 111      2  
## 2 112      1  
## 3 111      2  
## 4 113      1
```

summarize()

```
a %>%  
  group_by(c) %>% summarize(a1 = n_distinct(a))
```

```
## # A tibble: 3 × 2  
##   c      a1  
##   <chr> <int>  
## 1 111      2  
## 2 112      1  
## 3 113      1
```

Tips

You can use `()` for directly retrieve the variable. For example:

```
x <- 1:5  
x
```

```
## [1] 1 2 3 4 5
```

```
(x <- 1:5)
```

```
## [1] 1 2 3 4 5
```

Ungrouping

If you need to remove grouping, and return to operations on ungrouped data, use `ungroup()`: Notice: it's a common issue without removing grouping before another new grouping.

Without ungrouping

```
daily <- group_by(flights, year, month, day)
daily %>%
  group_by(month) %>%
  summarize(flights=n())
```

```
## # A tibble: 12 × 2
##   month flights
##   <int>   <int>
## 1     1    27004
## 2     2    24951
## 3     3    28834
## 4     4    28330
## 5     5    28796
## 6     6    28243
## 7     7    29425
## 8     8    29327
## 9     9    27574
## 10    10    28889
## 11    11    27268
## 12    12    28135
```

With ungrouping

```
daily %>%  
  group_by(month) %>%  
  ungroup() %>%  
  summarize(flights=n())
```

```
## # A tibble: 1 × 1  
##   flights  
##   <int>  
## 1  336776
```